

Building a High-Performance Mini SQLite-Based Database System

From Concept to Implementation

Database Systems Architecture

July 8, 2025

Table of Contents

- 1 Project Overview & Objectives
- 2 Core Architecture Overview
- 3 Core Components Design
- 4 Indexing Strategy
- 5 Query Processing Pipeline
- 6 Concurrency & Parallelization
- 7 Performance Optimizations
- 8 Implementation Roadmap
- 9 Testing & Benchmarking
- 10 Conclusion & Next Steps

Project Vision

Goal

Build a SQLite-inspired database system optimized for modern hardware constraints


Traditional SQLite Constraints (2000s):

- Limited RAM (64MB-1GB)
- Slow disk I/O
- Single-core CPUs
- Conservative memory usage


Modern Reality (2025):

- Abundant RAM (16-128GB)
- Fast NVMe SSDs
- Multi-core CPUs (8-32 cores)
- **Speed is everything**

System Architecture - High Level



Memory-First Architecture



Key Principles:

- Keep hot data in RAM
- Minimize disk I/O
- Optimize for cache locality
- Use memory-mapped files


Implementation:

- 80-90% RAM for buffer pool
- Async disk writes
- Prefetching algorithms

Storage Engine - Data Structures

Listing 1: Page Structure

```
1      class Page {  
2          static const size_t PAGE_SIZE = 4096;  
3  
4          struct Header {  
5              uint32_t page_id;  
6              uint16_t free_space;  
7              uint16_t slot_count;  
8              uint32_t checksum;  
9          };  
10  
11          Header header;  
12          std::vector<uint16_t> slot_directory;  
13          std::vector<uint8_t> data;  
14  
15      public:  
16          bool insert_record(const Record&  
17                          record);  
18          Record get_record(uint16_t slot_id);
```



Advantages:


- Cache-friendly layout
- Variable-length records

Buffer Pool Manager

Listing 2: Buffer Pool Implementation

```
1  class BufferPool {
2      struct BufferFrame {
3          Page* page;
4          uint32_t page_id;
5          bool is_dirty;
6          uint32_t pin_count;
7          std::chrono::steady_clock::time_point last_access;
8      };
9
10     std::unordered_map<uint32_t, BufferFrame*> page_table;
11     std::vector<BufferFrame> buffer_frames;
12     std::mutex buffer_mutex;
13
14     public:
15         Page* get_page(uint32_t page_id);
16         void unpin_page(uint32_t page_id, bool is_dirty);
17         void flush_all_pages();
18
19     private:
```

Buffer Pool - Memory Management Strategy




Skip List Index Implementation

Listing 3: Skip List Node

```

1      template<typename Key,
2              typename Value>
3      class SkipListNode {
4          Key key;
5          Value value;
6          std::vector<SkipListNode
7              *> forward;
8
9      public:
10         SkipListNode(Key k,
11                     Value v, int level
12                     )
13             : key(k), value(v),
14               forward(level
15                       + 1) {}
16
17         Key get_key() const {
18             return key; }

```

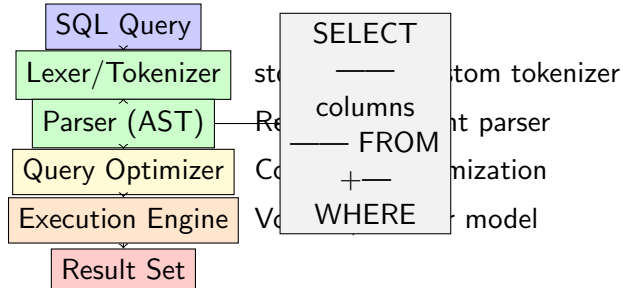


Hash Index for Equality Queries

Listing 4: Robin Hood Hash Table

```
1  template<typename Key, typename Value>
2  class RobinHoodHashTable {
3      struct Entry {
4          Key key;
5          Value value;
6          uint32_t hash;
7          uint32_t distance; // Distance from ideal position
8          bool is_occupied;
9      };
10
11      std::vector<Entry> table;
12      size_t capacity;
13      size_t size;
14
15  public:
16      bool insert(const Key& key, const Value& value);
17      bool find(const Key& key, Value& value);
18      bool remove(const Key& key);
19  }
```

Query Processing Flow



SQL Parser - AST Construction

Listing 5: Abstract Syntax Tree Nodes

```
1  class ASTNode {  
2  public:  
3      virtual ~ASTNode() = default;  
4      virtual void accept(ASTVisitor& visitor) = 0;  
5  };  
6  
7  class SelectStatement : public ASTNode {  
8      std::vector<std::unique_ptr<Expression>> select_list;  
9      std::unique_ptr<FromClause> from_clause;  
10     std::unique_ptr<WhereClause> where_clause;  
11     std::unique_ptr<OrderByClause> order_by_clause;  
12  
13     public:  
14         void accept(ASTVisitor& visitor) override {  
15             visitor.visit(*this);  
16         }  
17 };  
18  
19 class TableReference : public ASTNode {
```

Query Optimizer - Cost-Based Decisions

Listing 6: Query Optimization Framework

```
1  class QueryOptimizer {
2      struct Statistics {
3          size_t table_size;
4          size_t distinct_values;
5          double selectivity;
6          std::unordered_map<std::string, size_t> column_stats;
7      };
8
9      std::unordered_map<std::string, Statistics> table_stats;
10
11  public:
12      std::unique_ptr<ExecutionPlan> optimize(const ASTNode& query);
13
14  private:
15      double estimate_cost(const ExecutionPlan& plan);
16      std::vector<ExecutionPlan> generate_plans(const ASTNode& query);
17      ExecutionPlan select_best_plan(const std::vector<ExecutionPlan>& plans);
18
19      // Optimization rules
```

CPU Parallelization with OpenMP

Listing 7: Parallel Table Scan

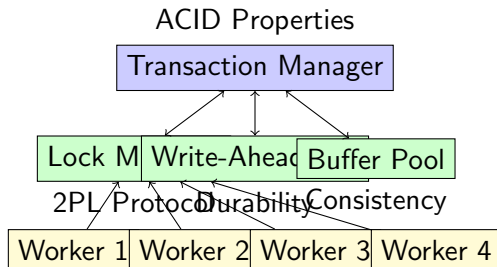
```
1  class TableScan {
2      std::vector<Page*> pages;
3      Predicate where_condition;
4
5  public:
6      std::vector<Record> scan_parallel() {
7          std::vector<Record> results;
8          std::mutex results_mutex;
9
10         #pragma omp parallel for
11         for (size_t i = 0; i < pages.size(); ++i) {
12             std::vector<Record> local_results;
13
14             // Scan page i
15             for (const auto& record : pages[i]->get_records()) {
16                 if (where_condition.evaluate(record)) {
17                     local_results.push_back(record);
18                 }
19             }
20         }
```

SIMD Vectorization for Batch Operations

Listing 8: Vectorized Comparison Operations

```
1  #include <immintrin.h>
2
3  class VectorizedOperations {
4  public:
5      // Compare 8 integers simultaneously
6      static std::vector<bool> compare_greater_than(
7          const std::vector<int32_t>& values,
8          int32_t threshold) {
9
10         std::vector<bool> results(values.size());
11         const __m256i threshold_vec = _mm256_set1_epi32(threshold);
12
13         size_t i = 0;
14         for (; i + 8 <= values.size(); i += 8) {
15             __m256i values_vec = _mm256_loadu_si256(
16                 reinterpret_cast<const __m256i*>(&values[i]));
17
18             __m256i cmp_result = _mm256_cmpgt_epi32(values_vec, threshold_vec);
```

Concurrency Control Architecture



Concurrency Strategy:

- **Two-Phase Locking (2PL):** Strict locking protocol
- **Write-Ahead Logging:** Ensure durability and recovery

Cache-Conscious Data Structures

Listing 9: Cache-Friendly Record Layout

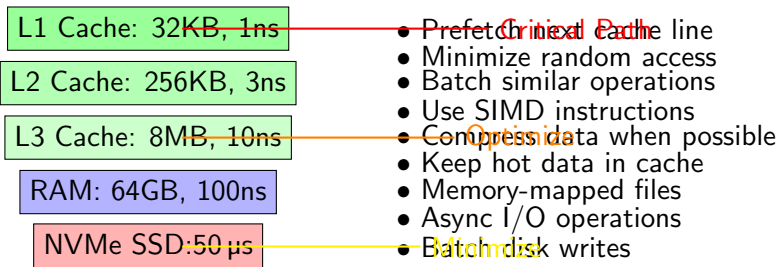
```
1 // Traditional approach - poor cache locality
2 struct Record_Bad {
3     std::string* name;    // Pointer to heap
4     int* age;             // Pointer to heap
5     double* salary;      // Pointer to heap
6 };
7
8 // Cache-friendly approach
9 struct Record_Good {
10     char name[64];        // Fixed-size, inline storage
11     int age;              // Value stored inline
12     double salary;        // Value stored inline
13
14     // Padding to align to cache line boundary
15     char padding[64 - sizeof(name) - sizeof(age) - sizeof(salary)];
16 } __attribute__((aligned(64)));
17
18 // Columnar storage for analytical queries
19 class ColumnStore {
```

Async I/O and Prefetching

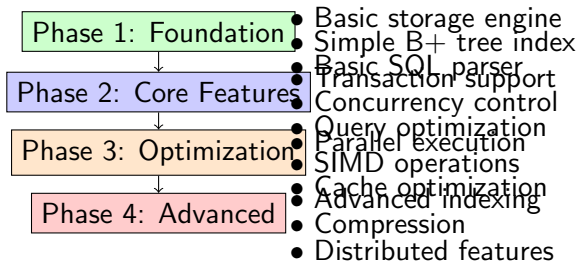
Listing 10: Asynchronous Disk Operations

```
1  #include <future>
2  #include <thread>
3
4  class AsyncDiskManager {
5      std::thread io_thread;
6      std::queue<std::packaged_task<void()>> io_queue;
7      std::mutex queue_mutex;
8      std::condition_variable cv;
9
10     public:
11         AsyncDiskManager() : io_thread(&AsyncDiskManager::io_worker, this) {}
12
13         std::future<Page*> read_page_async(uint32_t page_id) {
14             auto task = std::packaged_task<Page*>(){
15                 [this, page_id]() { return read_page_from_disk(page_id); };
16
17             auto future = task.get_future();
18
19             return future;
20         }
```

Memory Hierarchy Optimization



Development Phases



Phase 1: Foundation Implementation

Listing 11: Basic Storage Engine Skeleton

```
1 // Core classes to implement first
2 class StorageEngine {
3     std::unique_ptr<BufferPool> buffer_pool;
4     std::unique_ptr<DiskManager> disk_manager;
5
6     public:
7         bool create_table(const std::string& table_name,
8                           const TableSchema& schema);
9         bool insert_record(const std::string& table_name,
10                            const Record& record);
11         std::vector<Record> scan_table(const std::string& table_name);
12 };
13
14 class DiskManager {
15     std::fstream db_file;
16     std::mutex file_mutex;
17
18     public:
19     ReadNextPage(uint32_t page_id);
```

Phase 2: Adding SQL and Transactions

Listing 12: SQL Parser and Transaction Manager

```

1 // SQL Parser using recursive descent
2 class SQLParser {
3     std::regex select_pattern;
4     std::regex insert_pattern;
5     std::regex create_table_pattern;
6
7 public:
8     SQLParser() {
9         select_pattern = std::regex(R"(SELECT\s+(\.+)\s+FROM\s+(\w+) (?:\s+WHERE\s+(\.+)?)?)");
10        insert_pattern = std::regex(R"(INSERT\s+INTO\s+(\w+)\s+VALUES\s*\(((.+) \))");
11    }
12
13    std::unique_ptr<Statement> parse(const std::string& sql);
14 };
15
16 // Transaction Manager
17 class TransactionManager {
18     std::atomic<uint64_t> next_tid{1};

```

Phase 3: Performance Optimization

Listing 13: Parallel Query Execution

```
1 // Thread pool for parallel execution
2 class ThreadPool {
3     std::vector<std::thread> workers;
4     std::queue<std::function<void()>> tasks;
5     std::mutex queue_mutex;
6     std::condition_variable condition;
7     bool stop = false;
8
9 public:
10     ThreadPool(size_t num_threads = std::thread::hardware_concurrency());
11
12     template<typename F>
13     auto enqueue(F&& f) -> std::future<decltype(f())> {
14         auto task = std::make_shared<std::packaged_task<decltype(f())>>(
15             std::forward<F>(f));
16
17         auto future = task->get_future();
18         {
19             std::unique_lock<std::mutex> lock(queue_mutex);
```

Phase 4: Advanced Features

Listing 14: Compression and Advanced Indexing

```
1 // Data compression for storage efficiency
2 class CompressionManager {
3 public:
4     // Dictionary compression for strings
5     std::vector<uint8_t> compress_string_column(
6         const std::vector<std::string>& strings) {
7         std::unordered_map<std::string, uint32_t> dictionary;
8         std::vector<uint32_t> compressed_data;
9
10        uint32_t next_id = 0;
11        for (const auto& str : strings) {
12            if (dictionary.find(str) == dictionary.end()) {
13                dictionary[str] = next_id++;
14            }
15            compressed_data.push_back(dictionary[str]);
16        }
17
18        return serialize_compressed_data(dictionary, compressed_data);
19    }
```


Performance Testing Framework

Listing 15: Benchmark Suite Implementation

```
1  #include <chrono>
2  #include <random>
3
4  class DatabaseBenchmark {
5      std::unique_ptr<StorageEngine> engine;
6      std::mt19937 rng;
7
8  public:
9      // TPC-C style benchmark
10     void benchmark_oltp_workload() {
11         const size_t num_transactions = 10000;
12         const size_t num_threads = std::thread::hardware_concurrency();
13
14         auto start = std::chrono::high_resolution_clock::now();
15
16         std::vector<std::thread> workers;
17         for (size_t i = 0; i < num_threads; ++i) {
18             workers.emplace_back([this, num_transactions, num_threads, i]() {
19                 for (size_t j = 0; j < num_transactions; ++j) {
20                     // Perform a transaction
```

Performance Metrics & Targets

Operation	Target	Measurement	Comparison
Point Query	< 0.1ms	std::chrono	vs SQLite
Range Scan	1M rows/sec	Throughput	vs PostgreSQL
Insert	10K TPS	Transactions/sec	vs MySQL
Parallel Scan	4x speedup	vs serial	CPU cores
Memory Usage	80% buffer hit	Cache ratio	vs traditional

Query Time (ms)



Key Takeaways


Architecture Principles

- **Memory-first design:** Leverage abundant RAM for speed
- **CPU parallelization:** Use all available cores with OpenMP
- **Cache-conscious data structures:** Optimize for modern CPU hierarchy
- **Async I/O:** Overlap computation and disk operations

Implementation Strategy

- **Start simple:** Basic storage engine with file I/O
- **Iterate quickly:** Add features incrementally
- **Measure everything:** Benchmark each optimization

Development Timeline



Basic SQL Engine Transpilation Advanced Features
Deliverables:

- Working prototype
- Performance benchmarks
- Documentation
- Test suite

Questions & Discussion

Thank You!

Contact & Resources

- **Source code:** Will be available on GitHub
- **Documentation:** Comprehensive API documentation
- **Benchmarks:** Performance comparison results
- **Community:** Open source contribution guidelines